



General Info

Amazon Nova AI Hackathon

Where

AI Hire AI

We envision the 2nd Earth where AI comes to HIRE another AI

14-2 - 16/3 1 PM

AI Hire AI is a dual-sided multi-agent app that automates both candidate applications and recruiter workflows in a busy world. Candidates are guided to build a clean application packet, complete fast micro-screens, and submit to the best-fit roles with minimal effort. Recruiters get real-time routing, evidence-backed ranking, prompt cards, and verification risk signals during conversations. After applications, action agents push ATS-ready notes, tags, stage changes, scheduling, and follow-ups into the existing ATS so interviews are queued with no manual data entry.

Team:

Duy Dung, Trang Cao, Lam Anh, Que Anh, Duc Anh, Mai Tran, Tue Tran, Bao Tran, Nguyen Nguyen, Khoi Nguyen

UI UX Design: **Trang Cao**, Mai Tran, Que Anh, Nguyen Nguyen

AI Models: **Khoi Nguyen**, Bao Tran, Tue Tran, Nguyen Nguyen, Que Anh

Data: Tue Tran, Khoi Nguyen, **Nguyen Nguyen**, Duy Dung, Duc Anh

Market Research: **Que Anh**, Nguyen Nguyen, Duy Dung, Duc Anh

Fullstack: **Lam Anh**, **Bao Tran**, Khoi Nguyen, Mai Tran (Frontend), Nguyen Nguyen (Backend)

Demo Presenter: Trang Cao (edit), Nguyen Nguyen, Lam Anh

Note Taking: **Duc Anh**

Blog Writing: **Bao Tran**

DevPost Submission: **@everyone**

Github: <https://github.com/anhnamtruong/cafair>

AI Candidate: <https://volume-yodel-17191690.figma.site>

AI Recruiter: <https://brief-boho-07634752.figma.site>

Notion for Dev:

<https://www.notion.so/AI-Hire-AI-Milestone-1-310276cbaecd80bc9ea5f24b690f39cf>

Trang's Document:  [RecruitFlow_Competitive_Audit.docx](#)

Design info:  [AWS Hack 2026_Design Info](#)

Conceptual Design: <https://crate-raid-58610697.figma.site/>

Figma Design:

<https://www.figma.com/design/BpmwcCzUlePhtPyRbjNbPg/Amazon-Nova-Hack?>

Proposal: <https://www.overleaf.com/project/69924be376bcad0aa16c2fc3>

DevPost: <https://devpost.com/software/my-valentine>

<https://amazon-nova.devpost.com/>

Tentative Architecture Diagram:

<https://www.overleaf.com/8546367796tqqqcfqnrsmh#38fbdc>

Market Research:

Hiring Tech Trends:

<https://www.ibm.com/products/watsonx-orchestrate/ai-agent-for-hr>

<https://www.greenhouse.com/blog/greenhouse-2025-workforce-hiring-report>

<https://pages.technologyadvice.com/best-human-resources-software-nofp-lp>

OpenClaw: <https://github.com/openclaw/openclaw>

<https://youtu.be/OBNvaksQgec>

 **Build & Automate ANYTHING With AMAZON NOVA! Here is how**

Table fyi

Feature	Our Product	Handshake	Brazen	vFairs	Paradox	Workday ATS
Virtual fair (booths/sessions)	✓	✓	✓	✓	✓	✗
Pre-fair ranking + shortlist	✓	✗	✗	✗	✗	✗
Live routing lanes (Now/Screen/Redirect)	✓	✗	✗	✗	✗	✗
3-min micro-screen + scorecard	✓	✗	✗	✗	✓	✓
Resume vs JD match view (keyword highlights)	✓	✗	✗	✗	✗	✗
Verification signals (similarity/consistency)	✓	✗	✗	✗	✗	✗
One-click follow-ups (templates)	✓	✓	✓	✗	✓	✓
Scheduling automation (invite + slots)	✓	✓	✗	✗	✓	✓
ATS overlay/sync (works with existing ATS)	✓	✓	✓	✓	✓	✓
UI action agent for ATS workflows (Nova Act)	✓	✗	✗	✗	✗	✗
End-to-end to offer + Day 1 handoff	✓	✗	✗	✗	✓	✓

Legend:  Our Product  Supported  Not supported

Tech Stack

Layer	Tech (notes)
Experience / UI	Next.js 16 (TypeScript), Tailwind v4, shadcn/ui, tRPC/Hono, Clerk Auth, Figma
Backend / Agents	Node.js + TypeScript, agent modules (triage/verify/microscreen), AWS Lambda / App Runner, Nova 2 Lite + embeddings + Sonic
Data / Storage	S3 (artifacts), Postgres (Supabase) + RLS, Drizzle ORM + migrations, Redis (optional cache)
Automation / Integrations	Nova Act (UI workflows), ATS overlay sync (notes/tags/stage), scheduling + email, audit logs



Role Description

Role	People	Suggestions
UI/UX	Trang Cao, Mai Tran, Que Anh, Nguyen Nguyen	User flows; Figma screens; dashboard + candidate page; components for frontend.
AI Models	Khoi Nguyen, Bao Tran, Tue Tran, Nguyen Nguyen	Triage + verify + microscreen logic; rubrics/thresholds; evidence text; test cases.
Data	Tue Tran, Khoi Nguyen, Nguyen Nguyen	DB tables; demo dataset; sample candidates (good/ok/risky/bad); metrics numbers.
Market search	Re- Que Anh, Nguyen Nguyen	World's Innovation, Competitors, stalk other work on Devpost
Fullstack	Lam Anh, Bao Tran, Khoi Nguyen; Mai Tran (FE); Nguyen Nguyen (BE)	Build UI + APIs; connect agents; Nova calls; Nova Act ATS sync; deploy stable demo.
Demo	Trang Cao	3-min script; click path; rehearsals; backup plan; final video flow.
Notes	Duc Anh	Meeting notes; task list; blockers; daily status; capture links/screenshots.
Blog	Bao Tran	builder.aws post draft; architecture + impact; screenshots; final edit + submit.



Timeline

Timeline: 14-2 – 16/3 1 PM

14/2 – 20/2: Warm up, Finalize Idea

Today: 20/2

20/2 – 27/2: Milestone 1

28/2 – 7/3: Milestone 2

Spring Break(GMU, VT): 7/3 – 15/3 >> Grind

13/3: Product Ready

13-16/3: Video Demo, Write Blog, DevPost

16/3 1PM: Deadline to submit

*Xx: career fair product ready → incase have time:
future direction:*

Meeting Notes

- 21/2 Trang, Que Anh, Bao Tran, Nguyen Nguyen

Clean UI, clear UX, not too fancy UI, so user can focus on what's important in our product, highlight key features, consider colors, size, structure, animations...

Check accessibility score, hierarchy of colors <https://webaim.org>

Consider display quickly what and why is the reason user see and take action: click something. Build Critical User Journey properly

Analyze flow and give pros and cons

Computer Analysis: focus on user step-by-step flow, business standpoints, e.g. consider conversion rates for clicking on buttons, privacy concerns (data leakage)--> solved this by give user headups before clicking.

Recommend computer analysis to be done on iPad

Todo:

- ☐ Do Computer Analysis for UX UI (Que Anh, Mai, Nguyen)
- ☐ UPDATE Critical User Journey-CUJ (Nguyen, Trang)
- ☐ Setup requirements todos for technical: fullstack + list tech stacks needed (Bao)
- ☐ Check error accuracy for AI algorithm (Nguyen)

Tentative tech stacks fyi:

Layer	Tech (notes)
Experience / UI	Next.js 16 (TypeScript), Tailwind v4, shadcn/ui, tRPC/Hono, Clerk Auth, Figma
Backend / Agents	Node.js + TypeScript, agent modules (triage/verify/microscreen), AWS Lambda / App Runner, Nova 2 Lite + embeddings + Sonic
Data / Storage	S3 (artifacts), Postgres (Supabase) + RLS, Drizzle ORM + migrations, Redis (optional cache)
Automation / Integrations	Nova Act (UI workflows), ATS overlay sync (notes/tags/stage), scheduling + email, audit logs

- 20/2

Who was there? Nguyen, Trang, Mai, Nguyen, Lam Anh, Duc Anh

Todos:

Intro

Master Doc

Trang brainstorms idea and design, e.g. computer analysis task, CUJ,...

Workflow agreement: Figma → code

Roles, lead and what

Timeline

Sub-team meeting

Team name

Notes:

Team introduction

Discuss on where we really need AI, be realistic

Shift from recruiter only to AI Hire AI

Work on product for both recruiter and candidate, automate both process

Reduce scope for recruiter we had, so we spend time on candidate side app

Timeline walkthru

Decide sub-team next meeting for design team

Next meeting:

- Sub-meeting:

Design : Trang, Mai, Nguyen, Que Anh


Technical: Lam Anh, Bao, Tue, Khoi Nguyen, Duc Anh, Dung

- Next full meeting? TBD

Tech todo

- ☐ Setup github repo and invite
- ☐ Setup AWS and invite
- ☐ List tech stacks - aim for efficient and impactful stuff
- ☐ Duc Anh and Duy start to think of datasets

RECRUITER side



<https://brief-boho-07634752.figma.site>

Persona 1 — "The Overwhelmed Career Switcher"

Name: Marcus, 31 | Marketing Manager → Product Manager

Situation: Has the skills but struggles to reframe his experience for new roles. Applies to 20+ jobs with the same resume and hears nothing back.

Goals: Get seen by the right recruiters, understand what roles actually fit his background, and stop wasting time on bad-fit applications.

Pain Points: Doesn't know how to tailor applications, loses track of where he applied, dreads the "tell me about yourself" open-screen because he rambles.

What AI Hire AI solves for him: The structured application packet reframes his experience into product language. Micro-screens give him a script. Best-fit routing means he's only submitting where he actually has a shot.

Persona 2 — "The High-Volume Passive Candidate"

Name: Priya, 27 | Software Engineer, currently employed


Situation: Open to opportunities but not actively searching. Ignores LinkedIn messages because they're always irrelevant and require too much effort to respond.

Goals: Only engage with roles that are genuinely worth her time. Minimal friction — she has a full-time job.

Pain Points: Generic outreach, repetitive screening calls, having to re-explain her background every time.

What AI Hire AI solves for her: One structured packet does the heavy lifting. A fast micro-screen replaces the 30-minute phone screen. She only shows up when it's already a strong match.

CANDIDATE side



<https://volume-yodel-17191690.figma.site>

Persona 3 — "The Overloaded Agency Recruiter"

Name: Jordan, 34 | Technical Recruiter at a staffing firm

Situation: Managing 15+ open reqs simultaneously, spending hours on intake calls, ATS notes, and follow-up emails. Drowning in admin.

Goals: Move faster, fill more reqs, spend time actually talking to candidates instead of processing them.

Pain Points: Manual ATS data entry after every call, no standardized way to compare candidates across reqs, forgetting to follow up.

What AI Hire AI solves for him: Post-conversation agents auto-generate notes, tags, and stage updates. Prompt cards keep interviews consistent. Evidence-backed ranking lets him defend his shortlist instantly.

Persona 4 — "The Skeptical In-House TA Lead"

Name: Diana, 42 | Head of Talent Acquisition at a 500-person tech company

Situation: Under pressure to reduce time-to-hire. Has a small team. Wary of AI tools that promise everything and integrate with nothing.

Goals: Reduce time-to-hire, ensure quality and fairness in screening, keep her team lean without losing candidate experience quality.

Pain Points: Tools that don't sync with her ATS (Greenhouse/Lever), bias concerns with automated scoring, losing the "human feel" that her employer brand depends on.

What AI Hire AI solves for her: ATS-native outputs, transparent evidence-backed ranking (explainable, not a black box), and verification risk signals that protect against fraud without removing human judgment.

MVP

Personas



Candidate Side

Persona 1 — "The Overwhelmed Career Switcher"

Name: Marcus, 31 | Marketing Manager → Product Manager

Situation: Has the skills but struggles to reframe his experience for new roles. Applies to 20+ jobs with the same resume and hears nothing back.

Goals: Get seen by the right recruiters, understand what roles actually fit his background, and stop wasting time on bad-fit applications.

Pain Points: Doesn't know how to tailor applications, loses track of where he applied, dreads the "tell me about yourself" open-screen because he rambles.

What AI Hire AI solves for him: The structured application packet reframes his experience into product language. Micro-screens give him a script. Best-fit routing means he's only submitting where he actually has a shot.

Persona 2 — "The High-Volume Passive Candidate"

Name: Priya, 27 | Software Engineer, currently employed

Situation: Open to opportunities but not actively searching. Ignores LinkedIn messages because they're always irrelevant and require too much effort to respond.

Goals: Only engage with roles that are genuinely worth her time. Minimal friction — she has a full-time job.

Pain Points: Generic outreach, repetitive screening calls, having to re-explain her background every time.

What AI Hire AI solves for her: One structured packet does the heavy lifting. A fast micro-screen replaces the 30-minute phone screen. She only shows up when it's already a strong match.



Recruiter Side

Persona 3 — "The Overloaded Agency Recruiter"

Name: Jordan, 34 | Technical Recruiter at a staffing firm

Situation: Managing 15+ open reqs simultaneously, spending hours on intake calls, ATS notes, and follow-up emails. Drowning in admin.

Goals: Move faster, fill more reqs, spend time actually talking to candidates instead of processing them.

Pain Points: Manual ATS data entry after every call, no standardized way to compare candidates across reqs, forgetting to follow up.

What AI Hire AI solves for him: Post-conversation agents auto-generate notes, tags, and stage updates. Prompt cards keep interviews consistent. Evidence-backed ranking lets him defend his shortlist instantly.

Persona 4 — "The Skeptical In-House TA Lead"

Name: Diana, 42 | Head of Talent Acquisition at a 500-person tech company

Situation: Under pressure to reduce time-to-hire. Has a small team. Wary of AI tools that promise everything and integrate with nothing.

Goals: Reduce time-to-hire, ensure quality and fairness in screening, keep her team lean without losing candidate experience quality.

Pain Points: Tools that don't sync with her ATS (Greenhouse/Lever), bias concerns with automated scoring, losing the "human feel" that her employer brand depends on.

What AI Hire AI solves for her: ATS-native outputs, transparent evidence-backed ranking (explainable, not a black box), and verification risk signals that protect against fraud without removing human judgment.

The Persona Matrix at a Glance

Person	Role	Primary Job-to-be-Done	Biggest Fear
a			
Marcus	Career switcher	Get seen by the right roles	Being filtered out before a human sees him

Priya	Passive candidate	Engage with zero wasted time	Another irrelevant recruiter wasting her time
Jordan	Agency recruiter	Fill more reqs with less admin	Falling behind on pipeline hygiene
Diana	TA Lead	Reduce time-to-hire, stay compliant	AI bias liability + broken integrations

MVP Feature Prioritization

Tier 1 — Do First

High Value / Low-to-Medium Effort — These make the product work at all

#	Feature	Value	Effort	Why It's First
1	Structured Application Packet Builder	★★★★★	🔧🔧	The core candidate artifact everything else depends on
2	Job-to-Candidate Fit Scoring (rules-based)	★★★★★	🔧🔧	Enables routing without ML complexity upfront
3	Micro-Screen Question Flow (async, text-based)	★★★★★	🔧🔧	Replaces phone screen — highest friction point for both sides
4	Post-Conversation Note & Tag Auto-Generation	★★★★★	🔧🔧	Immediate time save for Jordan; biggest "wow" moment
5	Candidate Profile → ATS-Ready Output	★★★★★	🔧🔧	Non-negotiable for Diana; makes recruiters trust the platform

Tier 2 — Do Second

High Value / Medium Effort — These make the product good

#	Feature	Value	Effort	Why It's Second
---	---------	-------	--------	-----------------

6	Live Recruiter Prompt Cards (during conversation)	★★★★★	🔧🔧	High value but needs real-time infra; builds on packet data
7	Evidence-Backed Candidate Ranking Dashboard	★★★★★	🔧🔧	Diana's trust driver — needs fit scoring (Feature 2) first
8	Best-Fit Role Routing for Candidates	★★★★★	🔧🔧	Core to Marcus & Priya's value prop; needs job index
9	Automatic Stage Update Agent (post-conversation)	★★★★★	🔧🔧	Kills manual ATS hygiene for Jordan; needs post-conv pipeline
10	Follow-Up Email Auto-Draft Agent	★★★★★	🔧🔧	Quick win on top of note generation; same infra

Tier 3 — Do Third

Medium-High Value / Medium-High Effort — These make the product defensible

#	Feature	Value	Effort	Why It's Third
11	Verification Risk Signal Flags	★★★★★	🔧🔧	Important for Diana but needs signal data to be reliable
12	Scheduling Action Agent (calendar link + confirmation)	★★★★★	🔧🔧	High convenience; needs calendar API integrations
13	Recruiter Inbox / Candidate Queue View	★★★★★	🔧🔧	Needed for daily recruiter workflow — prioritize UX carefully
14	Candidate Application Status Tracker	★★★★★	🔧🔧	Reduces candidate anxiety; relatively simple to build
15	Multi-Req Management View (for agency recruiters)	★★★★★	🔧🔧	Jordan's power feature — complex but high retention driver

Tier 4 — Post-MVP

Lower urgency or too complex to build correctly in V1

#	Feature	Value	Effort	Why It Waits
16	ML-Based Fit Scoring (replace rules-based)	★★★★★	🔧🔧🔧	Needs data volume you won't have at launch
17	Bias Audit & Explainability Layer	★★★★★	🔧🔧🔧	Critical for Diana long-term; needs scoring infra first
18	Native ATS Integration (Greenhouse, Lever, Workday)	★★★★★	🔧🔧🔧	Huge value but deep API work; ship CSV export first
19	Candidate Skill Gap Coaching	★★★★	🔧🔧🔧	Nice-to-have for Marcus; doesn't drive recruiter adoption
20	Analytics & Recruiter Performance Dashboard	★★★★	🔧🔧🔧	Valuable for Diana but a distraction in MVP

Key Screens by Core Task

Candidate Side — 6 Core Screens

Screen 1 — Onboarding & Packet Builder

Core Task: Turn raw resume/LinkedIn into a structured application packet

This is the candidate's first real interaction with the platform. The screen needs to guide Marcus through uploading or importing his background, then intelligently parse and restructure it into a standardized packet — skills, experience, work style preferences, and role targets. The output isn't a resume. It's a structured profile the matching engine and recruiters actually consume.

Critical UI moments: Progress indicator so it doesn't feel endless, inline AI suggestions as fields populate, a "review before submit" summary card at the end.

Screen 2 — Role Match Feed

Core Task: Show candidates only roles worth applying to

Priya's entry point every time she returns. A ranked, filtered list of roles with a visible fit score and a one-sentence explanation of why each role matched. Not a job board — a curated queue. Each card should show role title, company, fit score, match reasoning, and a single CTA to begin the micro-screen.

Critical UI moments: Fit score with a plain-language reason ("Strong match — your Python and system design experience aligns with 4 of 5 requirements"), not just a number. Swipe-to-dismiss for roles that don't feel right.

Screen 3 — Micro-Screen Flow

Core Task: Complete a fast, async pre-screen for a specific role

5–7 adaptive questions specific to the role and the candidate's packet gaps. Text or short voice response. Feels more like a conversation than a form. This is where the candidate earns their ranking position on the recruiter side.

Critical UI moments: One question at a time to reduce overwhelm. Time estimates per question ("~30 seconds"). No going back — keeps responses authentic and reduces overthinking.

Screen 4 — Application Status Tracker

Core Task: Know where every application stands without emailing anyone

A clean pipeline view showing each submitted application, its current stage, and the last action taken. Marcus's main anxiety-reducer. Should surface next steps clearly — "Interview being scheduled," "Awaiting recruiter review," "Micro-screen completed."

Critical UI moments: Stage labels in plain language, not ATS jargon. A timestamp on last activity so candidates know they haven't been ghosted.

Screen 5 — Interview Prep Card

Core Task: Arrive prepared for the recruiter conversation

Generated automatically once an interview is scheduled. Surfaces the role requirements, the recruiter's likely focus areas based on the job description, and any gaps in the candidate's

packet they should be ready to address. Priya sees this and knows exactly what she's walking into.

Critical UI moments: Compact and skimmable — this is read 10 minutes before a call. "What they'll likely ask" section based on the role and the candidate's micro-screen answers.

Screen 6 — Notification & Action Center

Core Task: Never miss a next step

A lightweight hub for time-sensitive actions — complete your micro-screen, confirm your interview time, respond to a recruiter message. Keeps passive candidates like Priya engaged without requiring her to check in constantly.

Critical UI moments: Deep links straight into the relevant action. No noise — only actionable items, never informational clutter.

Recruiter Side — 7 Core Screens

Screen 1 — Candidate Queue / Inbox

Core Task: See ranked, routed candidates for open reqs in one place

Jordan's daily home screen. Candidates organized by req, sorted by fit score, flagged with risk signals and micro-screen completion status. Not a spreadsheet — an actionable queue. Each candidate card surfaces the most decision-relevant information without requiring a click.

Critical UI moments: Req switcher at the top so Jordan can move between 15 open roles without losing context. Color-coded priority signals. One-click to open the full candidate packet.

Screen 2 — Candidate Evidence Card

Core Task: Evaluate a candidate before and during a conversation

The full view of a single candidate — structured packet, fit score with evidence breakdown, micro-screen responses, verification risk flags, and previous interaction history. This is what Diana needs to defend a shortlist and what Jordan needs to decide who to call next.

Critical UI moments: Evidence breakdown shown as "Why this candidate fits" bullet points tied directly to job requirements — not a score in isolation. Risk flags displayed inline, not buried in a separate tab.

Screen 3 — Live Conversation Assistant

Core Task: Run a structured, consistent interview without manual effort

The screen open during an active recruiter-candidate conversation. Surfaces dynamic prompt cards — suggested questions based on the candidate's gaps, role requirements, and what hasn't been covered yet. Updates in real time as the conversation progresses.

Critical UI moments: Unobtrusive layout — the recruiter is talking to a person, not staring at a screen. Prompt cards dismissable with one tap. A live transcript running in the background (not front and center).

Screen 4 — Post-Conversation Action Review

Core Task: Approve and dispatch everything the agent generated after a call

After the conversation ends, the action agent surfaces a structured summary — auto-generated ATS notes, suggested tags, recommended stage update, draft follow-up email, and scheduling action if applicable. The recruiter reviews and approves in under 60 seconds instead of writing everything manually.

Critical UI moments: Each output editable inline before approving. A single "Approve All" button for speed, with individual overrides available. This screen is Jordan's biggest daily time save — it has to feel instant.

Screen 5 — Ranking & Shortlist Dashboard

Core Task: Compare candidates across a req and build a defensible shortlist

A side-by-side or stacked view of top candidates for a single role, each with their evidence-backed score, micro-screen highlights, and risk status. Diana uses this to review her team's pipeline and present to hiring managers.

Critical UI moments: Exportable as a structured summary for hiring manager review. Filters by score, risk flag, stage, and availability. "Why ranked here" tooltip on every score so the ranking is never a black box.

Screen 6 — Req Management Overview

Core Task: Monitor pipeline health across all open roles

The macro view for Jordan managing 15 reqs and Diana overseeing her team. Shows each open role, candidate volume by stage, average time-in-stage, and number of pending actions. Surfaces bottlenecks at a glance — "3 candidates awaiting follow-up in Software Engineer req since Tuesday."

Critical UI moments: Drill down to a single req with one click. Alerts for stale pipeline — candidates who haven't moved in X days. Keeps nothing hidden and nothing requiring manual counting.

Screen 7 — Verification Risk Flag Review

Core Task: Identify and act on fraud or inconsistency signals before they waste time

A dedicated or inline view that surfaces candidates with flagged inconsistencies — employment date gaps, skill claim mismatches against micro-screen responses, or résumé anomalies. Not an automatic disqualification — a signal for the recruiter to investigate.

Critical UI moments: Each flag explained in plain language ("Claimed 5 years Python experience; micro-screen responses suggest beginner-level familiarity"). A "mark reviewed" action so flags don't repeat. Integrated directly into the Evidence Card so it's never a separate workflow.

Screen Map Summary

Side	Screen	Primary Persona	Tier
Candidate	Packet Builder	Marcus	MVP Core
Candidate	Role Match Feed	Priya, Marcus	MVP Core
Candidate	Micro-Screen Flow	Marcus, Priya	MVP Core
Candidate	Application Status Tracker	Marcus	MVP Core
Candidate	Interview Prep Card	Priya	MVP — Sprint 2

Candidate	Notification Center	Priya	MVP — Sprint 2
Recruiter	Candidate Queue / Inbox	Jordan	MVP Core
Recruiter	Candidate Evidence Card	Jordan, Diana	MVP Core
Recruiter	Live Conversation Assistant	Jordan	MVP Core
Recruiter	Post-Conversation Action Review	Jordan	MVP Core
Recruiter	Ranking & Shortlist Dashboard	Diana	MVP — Sprint 2
Recruiter	Req Management Overview	Jordan, Diana	MVP — Sprint 2
Recruiter	Verification Risk Flag Review	Diana	MVP — Sprint 3

Recruiter Flow

PHASE 1 — Open Req & Candidate Intake

Step 1 — Recruiter opens req

👉 Vào **Screen 6 — Req Management Overview**

Jordan nhìn thấy:

- 15 open roles
- Pipeline health
- Bottlenecks
- Candidate volume by stage

 **Decision:**

- Ưu tiên req nào hôm nay?
- Role nào đang thiếu top candidates?

- Role nào đang bị stale?

Jordan click vào 1 req cụ thể.

PHASE 2 — Evaluate Candidates

Step 2 — View Ranked Queue

Screen 1 — Candidate Queue

Hệ thống đã:

- Rank candidates theo fit score
- Highlight micro-screen completed
- Show risk flags

Jordan không phải mở từng hồ sơ như spreadsheet.

 Decision:

- Ai gọi trước?
- Ai reject ngay?
- Ai cần review sâu?

Jordan click candidate top 1.

Step 3 — Deep Candidate Review

Screen 2 — Candidate Evidence Card

Jordan thấy:

- Fit score breakdown (WHY)
- Micro-screen answers
- Risk flags
- Interaction history

 Decision:

- Worth a call?
- Need clarification?
- Reject?

Nếu pass → Schedule call.

PHASE 3 — Live Conversation

Step 4 — Run Structured Interview

 **Screen 3 — Live Conversation Assistant**

Trong lúc call:

- Prompt cards gợi ý câu hỏi
- Real-time transcript
- Highlight missing signals

 System continuously updates:

- Confidence score
- Gaps covered

- Risk signal update

Call kết thúc.

PHASE 4 — Auto Actions

Step 5 — Post-Call Review

Screen 4 — Post-Conversation Action Review

System auto-generate:

- ATS notes
- Suggested tags
- Stage update recommendation
- Draft follow-up email
- Scheduling suggestion

Jordan:

- Edit inline nếu cần
- Click “Approve All”

 60 seconds instead of 10 minutes manual admin.

Candidate moves stage.

PHASE 5 — Shortlisting

Sau khi interview 5–10 người:

Step 6 — Build Shortlist

👉 Screen 5 — Ranking & Shortlist Dashboard

Jordan hoặc Diana:

- Compare candidates side-by-side
- See evidence-backed ranking
- Check risk flags
- Filter by availability

🎯 Output:

- Structured export for Hiring Manager
- “Why ranked here” explanation

Shortlist sent to Hiring Manager.

PHASE 6 — Hiring Manager Interview

(Platform may or may not handle this — depends scope)

If integrated:

- Hiring manager feedback added
- Score updated

Otherwise:

- Recruiter manually updates stage
-

PHASE 7 — Final Decision

Top candidate selected.

Recruiter:

- Moves to Offer stage
- Sends offer
- Updates ATS

If offer accepted:

 Status = Hired

PHASE 8 — Risk Control Layer (Parallel Flow)

Throughout the whole process:

 **Screen 7 — Verification Risk Flag Review**

Can trigger anytime:

- Resume inconsistency
- Skill mismatch
- Suspicious signals

Recruiter:

- Mark reviewed
- Investigate

- Escalate if needed
-

Condensed Flow (Ultra Clear)

1. Open Req → Screen 6
2. Review Ranked Queue → Screen 1
3. Open Candidate → Screen 2
4. Conduct Interview → Screen 3
5. Approve Auto Notes → Screen 4
6. Compare & Shortlist → Screen 5
7. Offer → Hire
8. Risk signals active throughout → Screen 7

Competitive Audit

Strategic Competitive Audit

Principal Product Strategy & UX Research Report

1. COMPETITOR IDENTIFICATION

Direct Competitors (AI-Native Recruiting Platforms)

Competitor	Core Claim
Paradox (Olivia)	Conversational AI for high-volume hiring
HireVue	AI-powered video interviewing + assessments
Beamery	Talent CRM + AI matching
Findem	Attribute-based AI talent search
Metaview	AI interview notes & summaries
Dover	AI-assisted sourcing + recruiter coordination

Indirect Competitors (Adjacent Threat Vectors)

Competitor	Why Relevant
LinkedIn Recruiter + AI features	Distribution moat + increasing automation
Greenhouse + AI integrations	ATS incumbency + workflow lock-in
Ashby	Modern ATS with built-in analytics targeting technical recruiting
Pallet	Community-native job routing
Teal / Simplify	Candidate-side automation (job tracking, auto-fill)
Rezi / Kickresume	Candidate document structuring
ChatGPT / Claude	Ad-hoc workflow replacement risk

2. COMPETITOR DEEP ANALYSIS

◆ Paradox (Olivia)

Core Positioning Conversational AI recruiter for high-volume, frontline hiring. Positions itself as the ATS replacement for enterprises running mass hiring (retail, logistics, QSR).

Primary User TA operations teams at enterprise scale (Walmart, McDonald's, Nestle). Not knowledge worker hiring.

Key UX Strengths

- SMS/WhatsApp-native interaction reduces drop-off dramatically in non-desk worker segments
- Scheduling automation is deeply integrated and genuinely eliminates coordinators
- Olivia persona creates pseudo-human warmth that outperforms cold system emails

UX Weaknesses

- Completely breaks down for nuanced, knowledge-worker evaluation
- Candidate experience is scripted and rigid — zero adaptability
- No candidate-side intelligence; purely recruiter-serving
- ATS notes still require manual review; summaries are shallow

Monetization Enterprise SaaS, seat-based + usage tiers. High ACV, long sales cycles, heavy IT dependency.

Emotional Design Approach Corporate reassurance. Designed to make TA directors feel safe and compliant, not delighted.

Trust Signals Fortune 500 case studies, compliance certifications, named enterprise logos prominently displayed.

◆ HireVue

Core Positioning Predictive hiring through structured video + game-based assessments. Sells "bias reduction" and structured evaluation consistency.

Primary User Mid-to-large enterprise TA teams running high-volume campus or standardized role hiring.

Key UX Strengths

- Strong assessment science narrative (sells to CHRO, not just TA)
- Async video normalizes the format for candidates in structured industries
- Deep psychometric data generates defensible ranking rationale

UX Weaknesses

- Candidate experience is anxiety-inducing — async video with AI scoring creates high abandonment
- Trust deficit: broadly perceived as dehumanizing; significant press backlash
- No live interaction intelligence; insight arrives post-hoc, not in-conversation
- No recruiter workflow automation post-assessment

Monetization Enterprise license + per-assessment fees. Upsells into analytics dashboards.

Emotional Design Approach Institutional authority. Designed to signal scientific rigor, which alienates candidates while reassuring buyers.

Trust Signals Published validation studies, EEOC compliance language, branded "fairness" reports.

Metaview

Core Positioning AI interview intelligence layer — records, transcribes, and summarizes structured interviews into recruiter-ready notes.

Primary User Technical recruiters and hiring managers at growth-stage tech companies who are drowning in manual note-taking.

Key UX Strengths

- Near-zero friction adoption (Zoom/Meet plugin, no workflow change required)
- Summary quality is genuinely high for structured interviews
- Positioned as assistant, not replacement — low internal resistance

UX Weaknesses

- Passive tool only — no pre-interview routing, ranking, or candidate intelligence
- No candidate-side product at all
- Insights are retrospective; nothing surfaces during the conversation
- No ATS write-back automation; still copy/paste for most teams

Monetization Seat-based SaaS. SMB-friendly pricing. Expansion through team invites.

Emotional Design Approach Quiet productivity tool. Minimal UI, focus on output quality. Appeals to time-starved IC recruiters.

Trust Signals Usage metrics ("X hours saved"), named tech company customers, privacy/GDPR compliance.

◆ Beamery

Core Positioning Talent Operating System — CRM + skills intelligence for workforce planning, not just active hiring.

Primary User Enterprise Talent Acquisition + Workforce Planning teams at global organizations.

Key UX Strengths

- Skills ontology is genuinely sophisticated — enables cross-role mobility mapping
- Strong talent pooling and nurture for passive candidates
- Integrates across HRIS + ATS ecosystems

UX Weaknesses

- Overwhelming complexity — product sprawl with no clear UX narrative
- Candidate experience is an afterthought; portal feels like a 2019 career page
- AI features feel bolted on, not architecturally native
- No real-time recruiter intelligence during live interactions

Monetization Enterprise license, professional services heavy. High implementation cost = high churn risk.

Emotional Design Approach Executive strategy tool. Designed to feel like a "platform" investment, not a tool recruiters love.

Trust Signals Analyst recognition (Gartner, Forrester), global enterprise logos, skills science narrative.

◆ Teal / Simplify (Candidate Side)

Core Positioning Job search productivity tools for candidates — auto-fill, tracking, resume optimization.

Primary User Active job seekers frustrated by repetitive application mechanics, primarily in tech and business roles.

Key UX Strengths

- Autofill dramatically reduces application friction
- Job tracker creates psychological sense of control for candidates
- Resume scoring creates engagement loops and return visits

UX Weaknesses

- No intelligent routing to best-fit roles — candidates still discover and decide manually
- No connection to recruiter side — entirely siloed product
- Application quality is not improved, only speed
- No verification, trust signals, or screening preparation

Monetization Freemium + premium subscription. Low ACV, high volume play. Resume tools drive conversion.

Emotional Design Approach Empowerment narrative. "You vs. the broken job market" framing. High emotional resonance with frustrated candidates.

Trust Signals Community testimonials, job offer success stories, transparent pricing.

Dover

Core Positioning AI-augmented recruiting service — blends software automation with human recruiting coordination for startups and scale-ups.

Primary User Series A–C startups without dedicated recruiting infrastructure.

Key UX Strengths

- Hybrid model (software + human) reduces adoption risk vs. pure AI
- Sourcing quality is high because human judgment remains in loop
- Startup-native UX language and pricing

UX Weaknesses

- Not truly autonomous — human dependency caps scalability
- No real-time recruiter tooling during live interviews
- Candidate experience is entirely managed by Dover, not portable
- ATS handoff is inconsistent

Monetization Monthly retainer + placement fees. Service company with software veneer.

Emotional Design Approach Trusted partner framing. "Your recruiting team, on demand." Reduces anxiety about AI replacement.

Trust Signals Named founder testimonials, transparent process documentation, recruiter bios.

3. SYNTHESIS

Structured Comparison Table

Dimension	Paradox	HireVue	Metaview	Beamery	Teal/Simplify	Dover	AI Hire AI
Candidate-Side Intelligence	✗	✗	✗	✗	✓ Partial	✗	✓ Full
Recruiter-Side Intelligence	✓ Partial	✓ Post-hoc	✓ Post-hoc	✓ CRM only	✗	✓ Human	✓ Real-Time
Live Conversation Intelligence	✗	✗	✗	✗	✗	✗	✓
Dual-Sided Design	✗	✗	✗	✗	✗	✗	✓
Post-Conversation Automation	✗	✗	✓ Partial	✗	✗	✗	✓ Full
ATS Write-Back	✓ Basic	✓ Basic	✗	✓	✗	✗	✓
Verification / Risk Signals	✗	✓ Partial	✗	✗	✗	✗	✓
Application Packet Structuring	✗	✗	✗	✗	✓ Partial	✗	✓
Best-Fit Routing	✗	✗	✗	✓ Talent pool	✗	✓ Human	✓
SMB / Startup Viable	✗	✗	✓	✗	✓	✓	TBD

Candidate UX Quality	Low	Low	N/A	Low	High	Managed	Target: High
----------------------	-----	-----	-----	-----	------	---------	--------------

Pattern Recognition Summary

Pattern 1: The Post-Hoc Intelligence Trap Every competitor with AI delivers insight *after* the high-value moment has passed. HireVue scores after video. Metaview summarizes after the interview. Paradox routes before but goes dark during. Nobody is inside the conversation in real time. This is the single largest unoccupied moment in the recruiting workflow.

Pattern 2: Dual-Sidedness Is Entirely Unaddressed Every product in this market picks a side — either recruiter tooling or candidate tooling — and ignores the other. This creates a structural asymmetry where both parties are operating with incomplete intelligence and the connection between them is still largely manual. No platform is simultaneously improving the candidate's preparation quality AND the recruiter's live decision quality.

Pattern 3: Post-Conversation Work Remains Manual Despite billions invested in recruiting automation, the minutes after a candidate conversation — note transcription, ATS updating, stage movement, scheduling follow-up, tagging, drafting outreach — are still largely manual. Metaview touches this but requires copy/paste; no platform has closed the loop into autonomous action agents.

Pattern 4: Verification and Trust Are Structural Gaps Background verification, credential validation, and fraud/misrepresentation risk signals are either absent or outsourced to third-party tools late in the process. No platform surfaces these signals proactively during the screening interaction itself.

Pattern 5: Candidate Experience Is a Product Afterthought With the exception of Teal/Simplify, candidate-facing UX is universally poor. Candidates are treated as data inputs, not users. This creates an enormous emotional gap — candidates are the most anxious, most motivated party in the transaction, and no platform is serving them with genuine intelligence and preparation support.

Market Saturation Areas (*Avoid / Differentiate Away From*)

- **Resume parsing and ranking** — Commoditized. Every ATS vendor, every AI wrapper does this. No moat.
- **Async video interviews** — HireVue owns enterprise, and candidate backlash is accelerating against the format.
- **Scheduling automation** — Paradox, Calendly, and Greenhouse integrations have largely solved this. Not a differentiator.

- **Job description optimization** — Dozens of point solutions. Low defensibility, low switching cost.
 - **High-volume / frontline hiring** — Paradox's moat is deep here. Competing is costly without specific operational infrastructure.
-

White Space Opportunities (*Strategic Vacancies*)

● **White Space 1: Real-Time Conversation Intelligence Layer** No competitor surfaces live guidance to recruiters *during* an active conversation. Prompt cards, risk flags, evidence-based follow-up questions in the moment — this is entirely unoccupied.

● **White Space 2: Candidate Preparation as a Product Surface** Structured application packets, micro-screen practice, role-fit clarity, and submission confidence — these don't exist as a designed product experience anywhere. Candidates navigate this alone. First platform to genuinely serve candidates with intelligence (not just autofill) earns loyalty and data.

● **White Space 3: Autonomous Post-Conversation Action Layer** Unstructured conversation → structured ATS reality is a painful manual gap that no platform has automated end-to-end. The recruiter experience ends the moment a call does; everything after is administrative debt. Eliminating this entirely is a 10–15 hour/week reclamation per recruiter.

● **White Space 4: Verification Risk Signals Embedded in Workflow** Surfacing credential anomalies, résumé inconsistencies, or identity verification flags *during* interaction — not after an offer — shortens costly mis-hire cycles. No platform does this in-context.

● **White Space 5: Dual-Sided Data Network** A platform that improves candidate quality (preparation, structure, honesty) and recruiter quality (live intelligence, decision support) simultaneously creates a data network effect neither side can access independently. Candidate signal quality improves recruiter outcomes; recruiter feedback loops improve candidate matching. This compounding is structurally unachievable for single-sided platforms.

STEP 4 — STRATEGIC TRANSLATION

Design Implications

1. Build for the Anxious Candidate First Candidate emotional state — uncertainty, fear of rejection, information asymmetry — is the product problem nobody is solving. AI Hire AI must design its candidate experience from an emotional journey map, not a feature list. Every

interaction should reduce anxiety and increase informed confidence. If candidates trust the platform, they submit better data, which makes the entire system smarter.

2. The Recruiter Interface Is a Live Decision Cockpit The recruiter-side UX should be conceived as a real-time command surface, not a dashboard. Information density should be tuned to the conversational moment — not all data at once, but the *right* data at the *right* conversational beat. This requires a fundamentally different IA than a standard ATS or CRM.

3. Post-Conversation Automation Must Be Zero-Touch Any friction in the action agent pipeline — any moment where a recruiter must review, correct, or initiate — undermines the value proposition. The design goal is a recruiter who ends a call and is functionally done with administrative work. ATS notes appear, tags are applied, follow-ups are drafted, stage is updated, next steps are scheduled. The cognitive exit from the conversation must be total.

4. Trust Architecture Must Be Visible to Both Sides Verification signals, application completeness scores, and consistency flags must be displayed with transparency UX — showing candidates what's visible about them, and showing recruiters what's been validated vs. self-reported. Opacity undermines trust; informed visibility builds it.

5. Micro-Screen Design Is a Core Differentiator The micro-screen format must be designed to feel like a conversation, not a test. Asynchronous, mobile-native, low time-commitment, with immediate intelligent feedback for candidates. This is the moment of greatest drop-off in the funnel; the UX here determines whether the platform achieves supply-side scale.

Feature Recommendations

Must-Have (Core Platform Integrity)

- Structured application packet builder with AI-guided field completion and role-fit scoring
- Real-time recruiter conversation HUD: candidate context, prompt cards, risk signals, open questions
- Post-conversation action agent: ATS note generation, tag application, stage update, scheduling trigger, follow-up draft
- Candidate micro-screen engine with adaptive question logic and instant structural feedback
- Verification signal layer: résumé consistency check, credential plausibility flags, surfaced in-session

Should-Have (Defensibility Layer)

- Candidate-side role-fit routing with explainable match rationale (not black-box scoring)
- Recruiter conversation analytics: talk ratio, topic coverage, response quality patterns over time

- ATS bidirectional sync (Greenhouse, Ashby, Lever, Workday) with field-mapping intelligence
- Candidate preparation intelligence: what to expect, what signals matter for this role, practice micro-screens

Could-Have (Network Effect Accelerators)

- Anonymized cross-platform candidate outcome data ("candidates with this profile convert at X% for roles like this")
- Recruiter performance benchmarks vs. cohort (time-to-decision, conversion rates, note quality)
- Candidate reputation portability: verified skills, past screen results, structured references — owned by candidate, shared selectively

Product Positioning Statement

For talent acquisition teams and serious job seekers who are exhausted by the manual, asymmetric, and time-destroying mechanics of modern recruiting, AI Hire AI is the first dual-sided recruiting intelligence platform that operates inside the conversation — not after it. Unlike point solutions that automate a single stage for a single user, AI Hire AI simultaneously elevates candidate preparation quality and recruiter decision quality in real time, then eliminates all post-conversation administrative work through autonomous action agents. The result is a recruiting process that is faster, fairer, more defensible, and structurally better for every party — without asking either side to change how they work.

Strategic Defensibility Summary

The most important thing AI Hire AI must understand about its competitive position: **every advantage is time-sensitive**. The white spaces above will not remain white. LinkedIn, Greenhouse, and Workday have distribution and data that could replicate individual features quickly. The defensible position is not any single feature — it is the **dual-sided data flywheel**: candidate quality signals improving recruiter outcomes, recruiter feedback improving candidate matching, compounding over every interaction. No single-sided platform can build this. It is the moat, and it must be architected from day one, not bolted on after product-market fit.

The second defensibility layer is **workflow depth over feature breadth**. Competitors that have tried to build platforms have created complexity without cohesion. AI Hire AI's design philosophy must be radical coherence — every feature serving the same moment, the same friction, the

same loop. Products that go deep into one workflow own it. Products that go wide lose to incumbents who already occupy the periphery.

Audit conducted at Principal Strategy level. Recommend follow-up primary research: 8–12 recruiter shadowing sessions focused on live conversation moments; 6–8 candidate journey mapping sessions focused on application anxiety points; ATS API capability audit across Greenhouse, Lever, Ashby, Workday for action agent integration depth assessment.